

Learning Objectives

- SQL Database Installation
- Databases **Terminology** Overview
- **Creating** – Databases, Schemas, Tables
- **Querying Data**
 - Basic SQL syntax
 - Filtering conditions
 - Alias
 - Searching Patterns
 - Duplicated rows
 - Sorting
 - Grouping
 - Data-type conversion...





Database,
DMBS, SQL

Database, DMBS, SQL

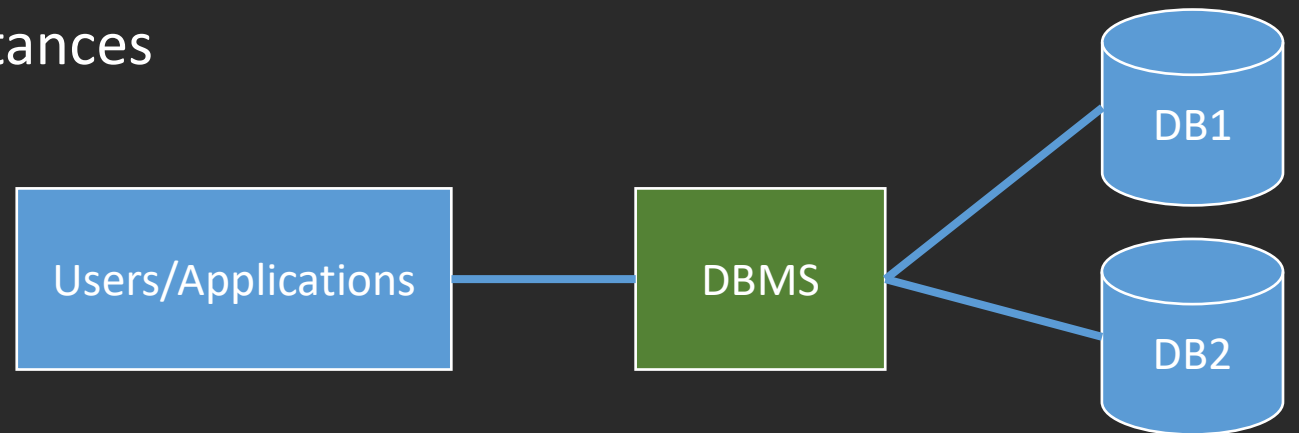
- What is a **Database**?

- An organized collection of data
- A logical container for storing data
- A **record** – one line combined from multiple attributes
 - Student : name, address, phone number.....
- Multiple **entities** that are connected to each other
 - Students
 - Teachers
 - Classes
 -

Database, DMBS, SQL

- **DBMS - Database Management System**

- Interface between users/Apps and databases
- Translate and process a query request
- Control access to the data
- Handle data integrity
- A unified layer to access the data
- Create and manage database instances
- **RDBMS** – Relational DBMS




Database, DMBS, SQL

- **SQL - Structured Query Language**

- Query, update, and modify data
- Stored and managed in **relational databases**
- The most popular language to query data
- Syntax - descriptive words (SELECT, FROM, WHERE....)
- Input and output are tables

- **CRUD**

- C – Create
- R – Read
- U – Update
- D – Delete



Tables, Columns, Rows

- **A Column**

- Related to specific **attribute**
- A pre-defined **data type**
 - String, numeric, datetime....
- Columns are also called **attributes**
- Data inside those columns are called **values**
- Attribute – Speed
 - Data type – numerical
 - Unit – KM/h

A diagram illustrating a grid structure. A 7x5 grid is shown with a green callout box pointing to the second column, labeled "Column". The grid consists of 7 rows and 5 columns. The second column is highlighted in green. The first, third, fourth, fifth, and sixth columns are gray. The fourth row is highlighted in blue.

Gray	Green	Gray	Gray	Gray
Gray	Green	Gray	Gray	Gray
Gray	Green	Gray	Gray	Gray
Blue	Green	Blue	Blue	Blue
Gray	Green	Gray	Gray	Gray
Gray	Green	Gray	Gray	Gray
Gray	Green	Gray	Gray	Gray

Tables, Columns, Rows

- **A Row**

- Collection of attributes related to an **entity**
 - Customers table: one row = one customer
- The number of rows is dynamic
 - 0 → Millions of records...
- Also called a **record/tuple**

Row

	Green	Grey	Grey	Grey
	Green	Grey	Grey	Grey
	Green	Grey	Grey	Grey
Blue	Blue	Blue	Blue	Blue
	Green	Grey	Grey	Grey
	Green	Grey	Grey	Grey
	Green	Grey	Grey	Grey

Primary and Foreign Keys

- **A Primary Key**

- Uniquely identify a **single row** in table
- Created from a single column/group of columns
 - Customers table → “Customer ID” as a primary key
- Defined when creating a table

Primary Key

	Column X			

Primary Key

	Column X	Column Y		

Primary and Foreign Keys

- A **Foreign Key**
 - Used to **navigate** from one table to another
 - **References** a primary key in another table
 - Created from a column or group of columns
 - Doesn't have to be unique

Orders

Attribute	Data Type	PK/FK
OrderID	Int	PK
OrderDate	Datetime	
Amount	Int	
CustomerID	Int	FK

Customers

Attribute	Data Type	PK/FK
ID	Int	PK
Name	varchar(30)	
Age	Int	
Phone	varchar(30)	



Relational Model, ER Diagram

- **Relational Model**

- **Entity**

- Distinguishable, unique, and distinct
 - All entities from the **same entity type** will be managed under the **same table**

- **Attribute**

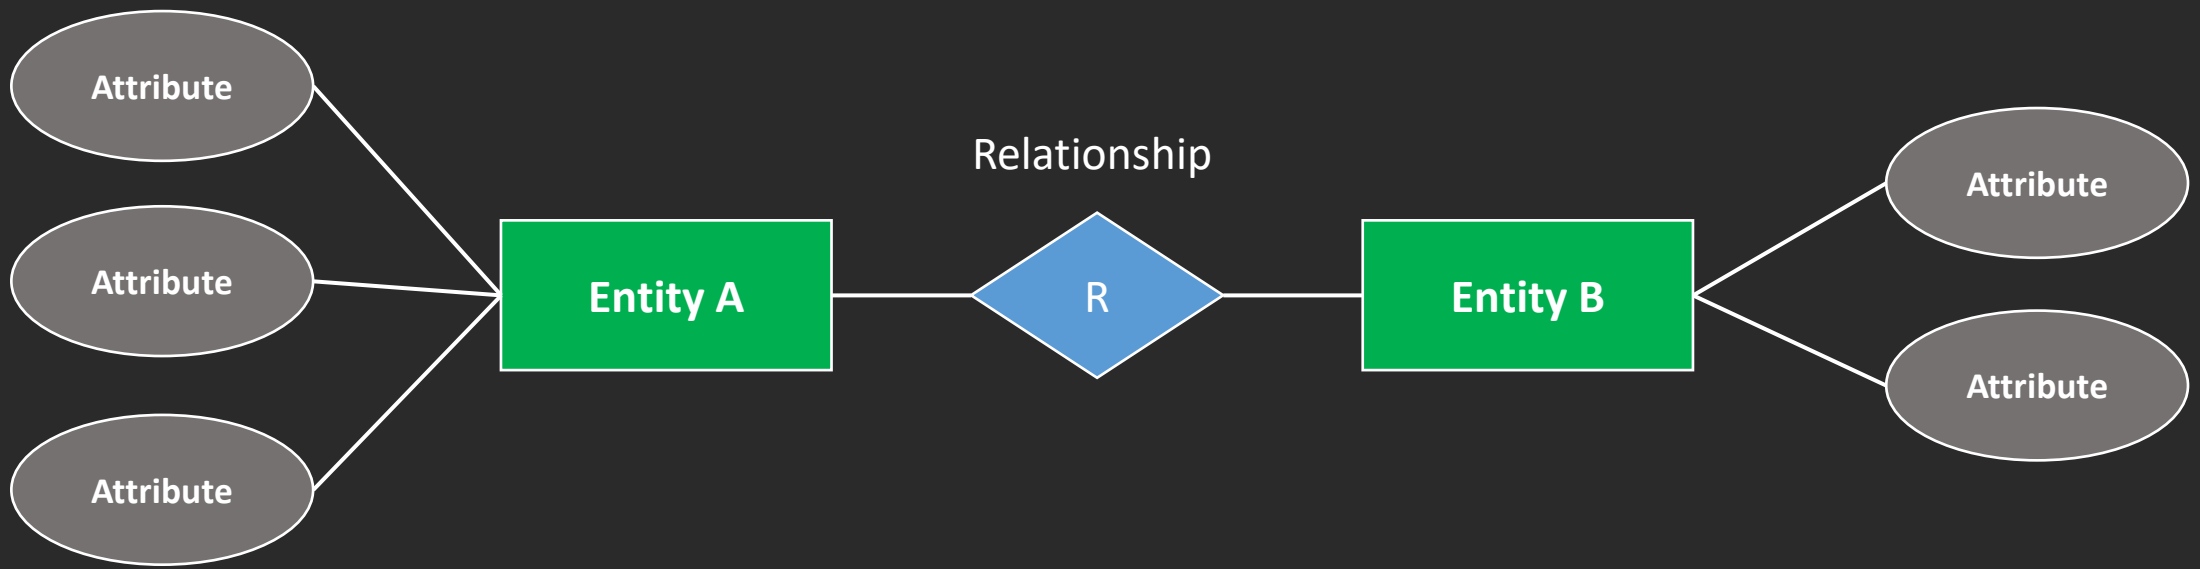
- A **characteristic** of an entity
 - Stored under a dedicated column (relevant data type)

- **Relationship**

- Describes **association** between different entities
 - One-to-many
 - Many-to-many
 - One-to-one

Relational Model, ER Diagram


- **ERD – Entity Relationship Diagram**
 - Illustrates how “entities” are related to each other
 - Model and design relational databases



Schema, Metadata, Data Dictionary

- Database **Logical Schema**

- The logical structure, constraints and associations of a **group** of **related tables**
- The “**blueprint**”- how the data must be stored and accessed
- Database instance → multiple schemas
 - **DB1.Schema1**
 - **DB1.Schema2**
- Each schema will have a **unique namespace**
- Qualify the table name
SCHEMA_NAME.TABLE_NAME

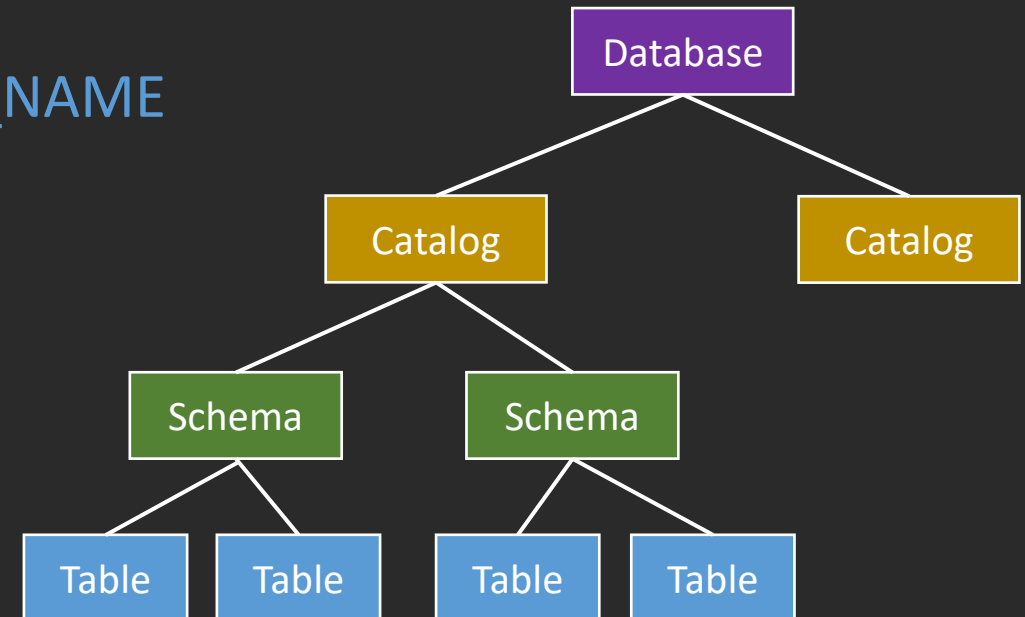


Schema, Metadata, Data Dictionary

- **Database Catalog**

- When using multiple schemas may not be enough
- Add another level to the overall hierarchy
- Qualify the table name

CATALOG_NAME.**SCHEMA_NAME**.**TABLE_NAME**



Metadata and Data Dictionary

- **Metadata**

- Data about the database **structure** and database **objects**
- For example –
 - List of tables in a schema
 - List of columns in a table
 - Column data type

- **Data Dictionary**

- Metadata information is managed in **dedicated tables** stored in the **data dictionary**

Null Values

- **What is a Null Value?**
 - “nothing” - the absence of any value
 - Used to mark missing data in columns
- **When inserting a Row (with missing data)**
 - Two Options
 1. Reject the row
 2. Insert new row and mark missing values as null
 - Based on the column definition
 - NULL value has a special marker

Indexes

- What is an **index**?
 - Rows are **NOT** saved in a **specific order** based on columns
 - DBMS can provide **sorted data**
 - List of students ordered by name
 - 100 records → 10,000,000 records
 - Requires a complete **table scan**
 - Inefficient and time-consuming process
 - An index - a mechanism to locate rows in a table


Indexes

- What is an **index**?

- An index is a **special support table** that is keeping rows in a specific order
 - Organizes the data by a **specified key**

Index Table on “Name”

Name	Id
Ava	8
Bob	55
David	7
Emma	14
Oliver	3
William	4



Students Table

Id	First Name	Last Name	Phone
7	David
3	Oliver
8	Ava
4	William
55	Bob
14	Emma

Partitions

- The **Big Data** Era

- More and More DATA
- HUGE challenge for databases
 - Collect, Store, Process...

- **Partitioning**

- A strategy to handle the growing amount of data
- Breaking a large table into smaller pieces called partitions

Partitions


- **Partitioning**

- Faster query - scan the required partition/s **instead** of the complete table
- Based on the **partition key**

IoT Sensors Table



Partitioning based on **YEAR**



Partitions

- **Optimized Queries**

- How a table was partitioned by the DB Admin?
- Include **filtering logic** on the **partitioning columns**
- Faster execution time!
- Cost-effective



Creating Databases, Schemas, and Tables



Retrieving Data with Queries

Query Data (SELECT)

Section	Purpose
SELECT	Determines which columns to include in the query's result set
FROM	Identifies the tables from which to retrieve data and how the tables should be joined
WHERE	Filters out unwanted data
GROUP BY	Used to group rows together by common column values
HAVING	Filters out unwanted groups
ORDER BY	Sorts the rows of the result set by one or more columns